

The zkEVM Architecture

Part V: Economics

Polygon zkEVM & Universitat Politècnica de Catalunya (UPC)

Marc Guzman-Albiol <marc.guzman.albiol@upc.edu>

Jose Luis Muñoz-Tapia <jose.luis.munoz@upc.edu>

Version: a26ddeaf521225eab1fd213e875f3d8e731766ea

January 18, 2024

User Fees

Basic Ethereum Fee Schema i

The basic fee schema to which Ethereum users are used works as follows.

The gas is a unit that accounts the resources used when processing a transaction.

At the time of sending a transaction, the user can decide two parameters:

1. **gasLimit:**

- It is the maximum amount of gas units that a user enables to be consumed by the transaction.

2. **gasPrice:**

- It refers to the amount of Wei a user is willing to pay per unit of gas for the transaction execution.
- In more detail, there is a market between users and network nodes such that if a user wants to prioritize his transaction, then he has to increase the **gasPrice**.

- At the **start of the transaction processing**, the following amount of Wei is subtracted from the source account balance:

$$\text{gasLimit} \cdot \text{gasPrice}.$$

- Then,
 - If $\text{gasUsed} > \text{gasLimit}$, the transaction is reverted.
 - Otherwise, the amount of Wei associated with the unused gas is refunded.
- The refunded amount of Wei that is added back to the source account is calculated as:

$$\text{gasLimit} \cdot \text{gasPrice} - \text{gasUsed} \cdot \text{gasPrice}.$$

Generic User Fee Strategy of Layer 2 Solutions

- In general, Layers 2 follow the fee strategy of charging an L2 gas price that is a percentage of the L1 `gasPrice`:

`L2GasPrice = L1GasPrice · L1GasPriceFactor.`

- For example:

`L1GasPrice = 20 Gwei`

`L1GasPriceFactor = 0.04 (4% of L1 gasPrice)`

`L2GasPrice = 20 Gwei · 0.04 = 0.8 Gwei`

- You can check the current fees at <https://l2fees.info>.

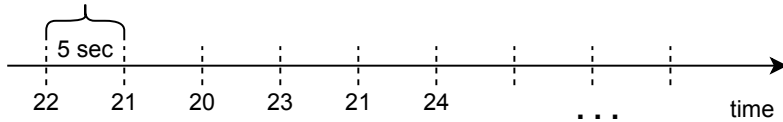
L2 User Fee Strategies are More Complex

However, this is not as easy as it may seem and there are additional aspects to consider:

- a) The **gasPrice** in L1 varies with time, so, how is this taken into account?
- b) Different **gasPrice** values in L1 can be used to prioritize transactions, how are these priorities managed by the L2 solution?
- c) The **gas/gasPrice** L1 schema may not be aligned with the actual resources spent by the L2 solution.

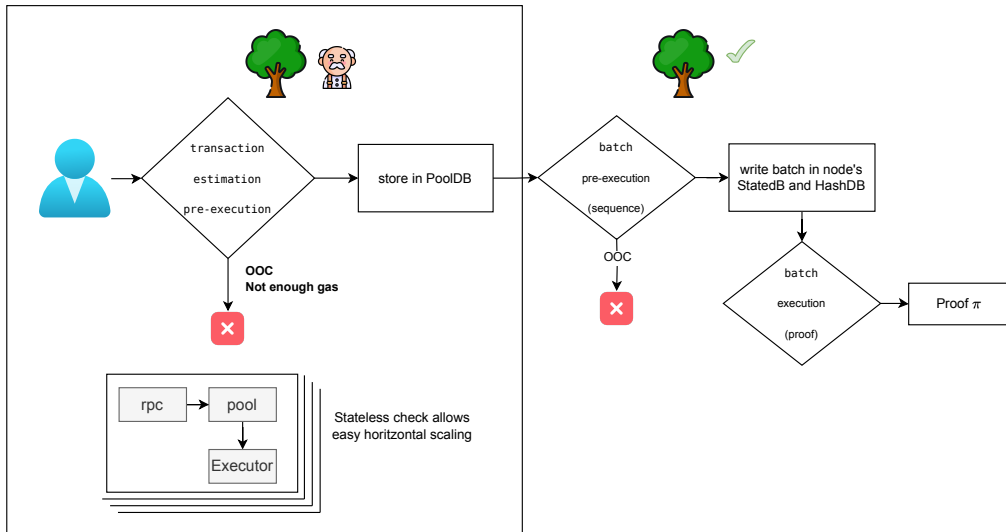
Obtaining L1 GasPrices

IntervalToRefreshGasPrices



In the example, we poll for the L1 gasPrice every 5 seconds and, as shown, gas prices vary with time.

RPC Transaction Pre-execution

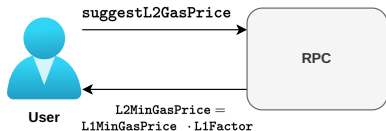


Sending a Transaction: User Experience

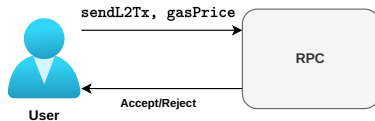
- We will operate in two steps: **gasPrice** suggestion and transaction sending.
- In the first step **A**, the user will ask via RPC call for a suggested **gasPrice** computed as

$$\text{L2GasPrice} = \text{L1GasPrice} \cdot \text{L1GasPriceFactor}$$

to sign its transaction with.



- In the second step **B**, the user sends the desired L2 transaction together with the **gasPrice**.

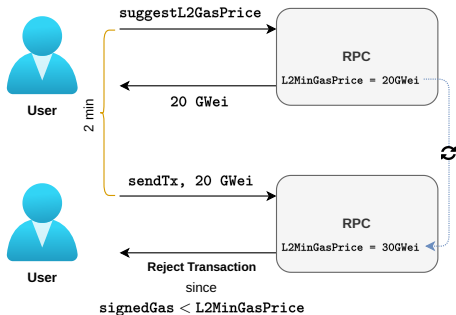


- If **gasPrice** provided by the user is less than the current **L2GasPrice**, the transaction is automatically **rejected**^a and not included into the pool (error `ErrGasPrice`).

^a Recall that the transaction can be rejected due to other checks explained before.

Sending a Transaction: Bad User Experience


- However, this previous schema is not a good UX.
- Between steps **A** and **B** there is an unbounded interval of time.
- Henceforth, the **L2GasPrice** from step **A** can be different from the one present in step **B**, leading to the following unwanted situation:



- Observe that, since the **L2GasPrice** has been refreshed, the transaction sent by the prover will be rejected even though it was signed with the exact suggested **gasPrice**.

Sending a Transaction: **MinAllowedPriceInterval**

- The solution is to allow transactions from users that have signed any **SignedGasPrice** that is above the minimum L2 gas price recorded during a period of time (called **MinAllowedPriceInterval**).
- This minimum is denoted as **L2MinGasPrice**.



We can configure the previous parameters in the Polygon zkEVM node:

```
1 [Pool]
2 ...
3 DefaultMinGasPriceAllowed = 0
4 MinAllowedGasPriceInterval = "5m"
5 PollMinAllowedGasPriceInterval = "1s"
6 IntervalToRefreshGasPrices = "5s"
7 ...
```

<https://github.com/0xPolygonHermes/zkevm-node/blob/develop/docs/config-file/node-config-doc.md#75-pooldb>

- **DefaultMinGasPriceAllowed:** It is the default min gas price to suggest.
- **MinAllowedGasPriceInterval:** It is the interval to look back of the suggested min gas price for a transaction.
- **PollMinAllowedGasPriceInterval:** It is the interval to poll L1 to find the suggested L2 min gas price.
- **IntervalToRefreshGasPrices:** It is the interval to refresh L2 gas prices.

When computing the L1 `gasPrice`, we can activate the `multigasprovider`:

```
1 [Ethereum]  
2 ...MultiGasProvider = false
```

When enabled, it allows using multiples sources for computing the L1 `gasPrice`.

L2 Gas Price Suggester


- However, with this particular design, the zkEVM endpoint that provides a suggestion for the gas price that the user has to sign with its transaction (which will be called L2 Gas Price Suggester) has a **big problem design**.
- Recall that the price of posting transactional data to L1 is charged to the zkEVM network to the **full L1 price**.
- Therefore, if we propose a gas price using **L1GasPriceFactor**, representing the measure of computational reduction in L2, there is a risk of running out of Wei reserves for posting data to L1.
- Consequently, we will recommend a slightly higher percentage of the gas price to the user, employing a **SuggesterFactor** of $0.15 \approx 4 \cdot \text{L1GasPriceFactor}$:

$$\text{GasPriceSuggested} = \text{L1GasPrice} \cdot \text{SuggestedFactor}.$$

Gas Price Suggester


```
1 [L2GasPriceSuggester]
2 Type = "follower"
3 UpdatePeriod = "3s"
4 Factor = 0.12
5 DefaultGasPriceWei = 0
6 MaxGasPriceWei = 0
7 CleanHistoryPeriod = "1h"
8 CleanHistoryTimeRetention = "5m"
```

Numeric Example: Minimum Gas Price i



- Observe that, when the user queries the suggested gas price through the RPC, the network responds with the current suggested gas price computed as $0.04 \cdot 19$, which is the current L1 gas price updated every 5 seconds.

Numeric Example: Minimum Gas Price ii



- However, at the time of sending the transaction, the RPC will only accept the transaction if **GasPriceSigned** is strictly higher than the minimum suggested gas price from 5 minutes ago (highlighted in **bold** in the figure), which in this instance is $19 \cdot 0.04 = 2.85$.
- In order to get his transaction accepted, the user sets the gas price of the transaction to $\text{GasPriceSigned} = 3.3 > 2.85 = \text{L2MinGasPrice}$.

L1/L2 Costs Issues

- Gas in Ethereum accounts the resources used by the transaction.
- In particular, it takes into account:
 - **Data availability** (the transaction bytes).
 - **Processing resources**, like CPU, Memory and Storage.
- Ethereum users are used to prioritize their transaction by increasing **gasPrice**.
- A big issue is that **there can be operations that consume low gas in L1 but that represent a major cost for L2**.
- The data availability costs are fixed once the transaction is known and they are directly proportional to L1 data availability costs.
- However, L2 execution is variable (because it depends on the state) and usually offers a smaller cost per gas.
- Henceforth, L2 transactions having high data availability costs and small execution costs are **highly problematic** in our pricing schema.

L1/L2 Costs Strategies

- Recall that the Ethereum fee is computed as $\text{gasUsed} \cdot \text{gasPrice}$, giving us two ways of solving the misalignment problem:

(A) Arbitrum Approach. Increase **gasUsed**.

- This approach is based on changing the gas schema to increase the Gas costs for data availability.
- This strategy is a relatively simple to implement and easy to understand but **it changes the Ethereum protocol**.
- An L1 Ethereum transaction may execute different when compared to the same transaction executed in L2.

(B) Effective Gas Price Approach. Increase **gasPrice**.

- If we do not want to modify the Gas, we have to increase **gasPrice** in order to cover the costs.
- Unlike the previous approach, this does not change the Ethereum specifications.
- However, it is complex to achieve a fair **gasPrice**.
- Moreover, we have to take into account that L2 users should be able to prioritize its transactions also increasing **gasPrice**, as they are used to.
- This is actually our approach.

Effective Gas Price Overview i

- The user signs a relatively high gas price at the time of sending the L2 transaction.
- Later on, by pre-executing the sent transaction, the **sequencer** establishes a fair **gasPrice** according to the amount of resources used.
- To do so, the **sequencer** provides a single byte **EffectivePercentageByte** $\in \{0, 1, \dots, 255\}$ (1 Byte), which will be used to compute a ratio called **effectivePercentage**

$$\text{EffectivePercentage} = \frac{1 + \text{EffectivePercentageByte}}{256}.$$

- The **effectivePercentage** will be used in order to compute the factor of the signed transaction's **gasPrice** which should be charged to the user:

$$\text{TxGasPrice} = \left\lfloor \text{GasPriceSigned} \cdot \frac{1 + \text{EffectivePercentageByte}}{256} \right\rfloor.$$

Effective Gas Price Overview ii

- For example, setting an `EffectivePercentageByte` of $255 = 0xFF$ would mean that the user would pay the totality of the `gasPrice` signed when sending the transaction:

$$\text{TxGasPrice} = \text{GasPriceSigned}.$$

- In contrast, setting `EffectivePercentageByte` to 127 would reduce the `gasPrice` signed by the user to the half:

$$\text{TxGasPrice} = \frac{\text{GasPriceSigned}}{2}.$$

- Observe that, in this schema, users **must trust the sequencer**.
- As having `EffectivePercentage` implies having `EffectivePercentageByte`, and vice versa, we will abuse of notation and use them interchangeably as `EffectivePercentage`.

About the **EffectivePercentage** computation

- We could account the pricing resources by means of the number of **consumed counters** present in our proving system.
- Nevertheless, comprehending this can be challenging for users, and it is crucial to prioritize a positive user experience in this specific aspect.
- Moreover, stating the efficiency through counters is not intuitive for users at the time of prioritizing their transactions.
- Henceforth, our actual goal is to compute **EffectivePercentage** only by using **Gas** and prioritizing users transactions by means of using **gasPrice**.

Introduction of the **BreakEvenGasPrice**

- Our goal as service providers is to **not accept transactions in which we loose money**.
- In order to achieve this, we will calculate the **BreakEvenGasPrice**, considering a secure threshold to avoid losses in the event of unexpected issues.
- As explained before, we will split the computation in two to take into account differently costs associated with data availability and costs associated with used Gas.

BreakEvenGasPrice: Costs Associated with Data Availability i

- Costs associated with Data Availability will be computed as

$$\text{DataCost} \cdot \text{L1GasPrice},$$

where **dataCost** is the cost in Gas for data in L1.

- In the Ethereum ecosystem, the cost of data varies depending on whether it involves zero bytes or non-zero bytes

$$\text{NonZeroByteGasCost} = 16, \quad \text{ZeroByteGasCost} = 4$$

- In particular, **non-zero bytes** cost 16 Gas meanwhile **zero bytes** 4 Gas.

BreakEvenGasPrice: Costs Associated with Data Availability ii

- Also recall that, when computing non-zero bytes cost we should take into account some constant data¹, always appearing in a transaction:
 - The **signature**, consisting on 65 bytes.
 - The **EffectivePercentageBytesLength**, consisting on 1 bytes related to the RLP-encoded fields length.
- This results in a total of 66 constantly present bytes.
- Taking all in consideration, **DataCost** can be computed as:

$$(TxConstBytes + TxNonZeroBytes) \cdot NonZeroByteGasCost + TxZeroBytes \cdot ZerByteGasCost,$$

where **TxZeroBytes** (resp. **TxNonZeroBytes**) represents the count of zero bytes (resp. non-zero bytes) in the raw transaction sent by the user.

¹This data can obtain zero bytes, but to optimize a little bit the processing we count them all of them as non-zero bytes.

BreakEvenGasPrice: Computational Costs

- For the computational cost, we will simply use the following formula:

$$\text{GasUsed} \cdot \text{L2GasPrice},$$

where recall that we can obtain `L2GasPrice` by multiplying `L1GasPrice` by chosen factor less than 1:

$$\text{L2GasPrice} = \text{L1GasPrice} \cdot \text{L1GasPriceFactor}.$$

- In particular, we will choose a factor of 0.04

$$\text{L1GasPriceFactor} = 0.04.$$

- Observe that, unlike data costs, in order to compute computational costs we will need to **execute** the transaction.

BreakEvenGasPrice Formula

- Now, combining both **data** and **computational** costs, we will refer to it as **TotalTxPrice**:

$$\text{TotalTxPrice} = \text{DataCost} \cdot \text{L1GasPrice} + \text{GasUsed} \cdot \text{L1GasPrice} \cdot \text{L1GasPriceFactor}.$$

- We can compute **BreakEvenGasPrice** as the following ratio:

$$\text{BreakEvenGasPrice} = \frac{\text{TotalTxPrice}}{\text{GasUsed}}.$$

- This calculation helps to establish the gas price at which the total transaction cost is covered.
- Additionally, we incorporate a factor **NetProfit** ≥ 1 that allows us to achieve a slight profit margin:


$$\text{BreakEvenGasPrice} = \frac{\text{TotalTxPrice}}{\text{GasUsed}} \cdot \text{NetProfit}.$$

- Observe that we still need to introduce here **gasPrice** prioritization.

Numeric Example: BreakEvenGasPrice i

- Recall the example proposed before, where the user ended up by setting `GasPriceSigned` to 2.85.
- Suppose the user sends a transaction having:
 - 200 non-zero bytes, including the constant ones.
 - 100 zero bytes.
- Moreover, at the time of pre-executing the transaction (without getting an **OOB** error), 60,000 Gas is consumed (recall that, since we are using a *wrong* state root, this gas is only an estimation).

Numeric Example: BreakEvenGasPrice ii



- Hence, the total transaction cost is of

$$(200 \cdot 16 + 100 \cdot 4) \cdot 21 + 60,000 \cdot 21 \cdot 0.04 = 126,000 \text{ GWei.}$$

- Observe that 21 is the **L1GasPrice** at the time of sending the transaction.

Numeric Example: **BreakEvenGasPrice** iii

- Now, we are able to compute the **BreakEvenGasPrice** as

$$\text{BreakEvenGasPrice} = \frac{\text{TotalTxPrice}}{\text{GasUsed}} = \frac{126,000 \text{ GWei}}{60,000 \text{ Gas}} \cdot 1,2 = 2.52 \text{ GWei/Gas.}$$

- Observe that we have introduced a **NetProfit** value of 1.2, indicating a target of a 20% gain in this process.
- At a first glance, we might conclude acceptance since **GasPriceSigned** = 3.3 > 2.52 but, recall that this is only an estimation, gas consumed with the correct state root can differ.
- Therefore, we introduce a **BreakEvenFactor** of 30% to account for estimation uncertainties:

$$\text{GasPriceSigned} = 3.3 > 3.276 = 2.52 \cdot 1.3 = \text{BreakEvenGasPrice} \cdot \text{BreakEvenFactor}.$$

- Consequently, we decide to **accept the transaction**.

Numeric Example: **BreakEvenFactor** i

- Imagine we disable the **BreakEvenFactor** setting it to 1.
- Our original transaction's pre-execution consumed 60k Gas, **GasUsedRPC** = 60k.
- However, imagine that the correct execution at the time of sequencing consumes 35k Gas.
- If we recompute **BreakEvenGasPrice** using this updated used gas, we get 3.6 GWei/Gas, which is way higher than the original one.
- That means that, we should have charged the user with a higher gas price in order to cover the whole transaction cost, which now is of 105,000 GWei.
- But, since we are accepting all the transactions signing more than 2.85 of gas price, we do not have margin to increase more.

Numeric Example: **BreakEvenFactor** ii

- In the worst case we are loosing

$$105,000 - 35,000 \cdot 2.85 = 5,250 \text{ GWei.}$$

- Introducing **BreakEvenFactor** we are limiting the accepted transactions to the ones having

$$\text{GasPriceSigned} \geq 3.27,$$

in order to compensate such losses.

- In this case, we have the flexibility to avoid losses and adjust both user and our benefits since

$$105,000 - 35,000 \cdot 3.27 < 0.$$

- Prioritization of transactions in Ethereum is determined by the signed **gasPrice**; higher values result in higher priority.
- To implement this, consider that users are only aware of two **gasPrice** values: the one signed with the transaction, called **GasPriceSigned**, and the one obtained from the RPC, that we will call **GasPriceSuggested**.

Introducing Priority ii

- In the case that `gasPriceSigned > gasPriceSuggested`, we establish a priority ratio as follows:

$$\text{PriorityRatio} = \frac{\text{GasPriceSigned}}{\text{GasPriceSuggested}} - 1.$$

- If `GasPriceSigned ≤ GasPriceSuggested`, the user has chosen not to prioritize its transaction (and maybe we can reject the transaction due to low gas price).
- In this case, we establish a priority ratio to be 0.
- The `EffectiveGasPrice` will be computed as:

$$\text{EffectiveGasPrice} = \text{BreakEvenGasPrice} \cdot (1 + \text{PriorityRatio}).$$

Numeric Example: **EffectiveGasPrice**

- Recall that, in the example, we were signing a gas price of 3.3 at the time of sending the transaction.
- However, the suggested gas price was of

$$\text{GasPriceSuggested} = 2.85.$$

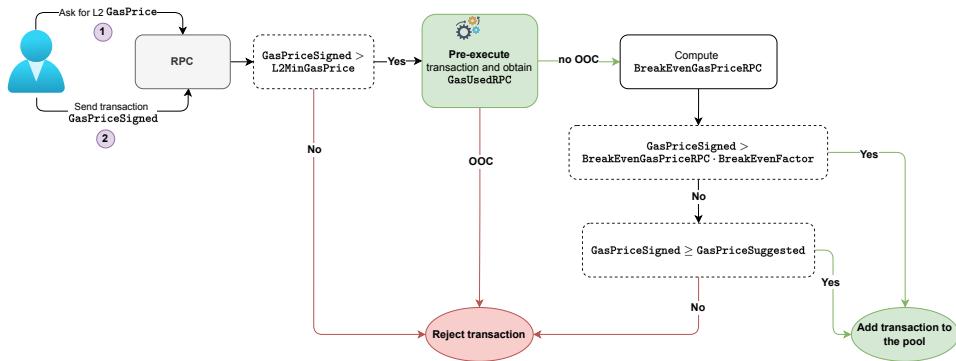
- The difference between both is taken into account in the priority ratio:

$$\text{PriorityRatio} = \frac{3.3}{2.85} - 1 = 0.16.$$

- Henceforth, the estimated **EffectiveGasPrice** in the RPC will be

$$\text{EffectiveGasPrice} = 2.52 \cdot (1 + 0.16) = 2.923.$$

gasPrice Flows: RPC i



1. The user asks to the RPC for a suggested L2 **GasPrice**.
2. The users sends the transaction together with a selected **GasPriceSigned**.
3. The RPC pre-executes the transaction (**important**, using a wrong state root) and uses the gas used **GasUsedRPC** in order to compute the **BreakEvenGasPriceRPC**.
4. We have two cases here:
 - If the transaction pre-execution runs out of counters (**OCC** error), we immediately reject the transaction.
 - If not, the RPC computes the **BreakEvenGasPriceRPC** and we continue the flow.

5. Now, we have two options:

- If $\text{GasPriceSigned} > \text{BreakEvenGasPriceRPC} \cdot \text{BreakEvenFactor}$, we immediately accept the transaction, storing it in the pool.
- **BreakEvenFactor** (which is equal to 1.3) is introduced to provide a wider safeness threshold.
- Otherwise $\text{GasPriceSigned} \leq \text{BreakEvenGasPriceRPC} \cdot \text{BreakEvenFactor}$, we are in dangerous zone because we may be facing losings.

6. In the bad path, we allow two options:

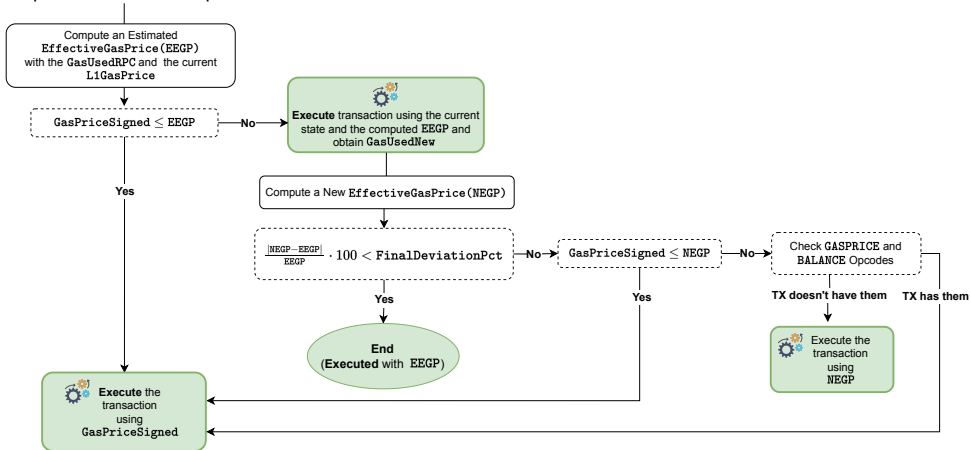
- If $\text{GasPriceSigned} \geq \text{GasPriceSuggested}$, we take the risk and we introduce the transaction into the pool.
- Otherwise $\text{GasPriceSigned} < \text{GasPriceSuggested}$, we immediately reject the transaction because its **highly probable** that we face losings.

gasPrice Flows: RPC. Some Considerations

- It is important to remark that, **once a transaction is included into the pool, we should actually include it into a block.**
- Hence, if something goes bad in later steps and the processing consumes far more gas than expected, we will **lose money having no possibility to overcome that situation.**
- On the contrary, if the process goes well and the processing consumes less gas than expected, we can reward the user by modifying the previously introduced **effectivePercentage**.
- Also observe that all the transactions stored in the **Pool** should be ordered from larger to lower priorities (or, equivalently, **effectiveGasPrice**).

gasPrice Flows: Sequencer i

Sequence a transaction from the pool



1. The sequencer takes a transaction from the **Pool** and recomputes the estimated **EffectiveGasPrice** (which we will call **EEGP**) using the **GasUsedRPC** stored by the RPC and the current **L1GasPrice**.
2. At this point, we have two options:
 - If $\text{GasPriceSigned} \leq \text{EEGP}$, there is a risk of loss.
 - In such cases, the user is charged the full **GasPriceSigned** and we end up the flow.
 - Conversely, if $\text{GasPriceSigned} > \text{EEGP}$, there is room for further adjustment of the user's gas price.
3. In this case, we recompute a new **EffectiveGasPrice** (which we will call **NEGP**) with the execution-related data compute from the **correct state root**.

4. We have two paths:

- If the difference between EEGP and NEGP is higher than some a parameter `FinalDeviationPct` (which is 10 in the actual configuration):

$$\frac{|\text{NEGP} - \text{EEGP}|}{\text{EEGP}} \cdot 100 < \text{FinalDeviationPct},$$

we end up the flow just to avoid re-executions and save execution resources.

- On the contrary, if the difference equals or exceeds the deviation parameter, there is a big difference between executions and we may better adjust gas price.

5. In the later case, two options arise:

- If `GasPriceSigned` \leq `NEGP` there is again a risk of loss.
- In such cases, the user is charged the full `GasPriceSigned` and we end up the flow.
- Otherwise, if `GasPriceSigned` $>$ `NEGP`, means that we have margin to adjust the gas price.
- However, we want to **save executions**, leading us to end up the process using a trick explained below.

6. We check if the transaction processing includes the two opcodes that use the gas price:
 - The **GASPRICE** opcode.
 - The **BALANCE** opcode from the source address.
7. If it is the case, to save one execution, we simply execute the transaction using the full **GasPriceSigned** to ensure we minimize potential losses and we end up the flow, as before.
8. If not, and with the intention of optimizing an execution while making a slight adjustment to the gasPrice, we proceed by executing the transaction using the **NEGP**.

Pool Effective Gas Price

```
1  ...  
2  [Pool.EffectiveGasPrice]  
3  Enabled = false  
4  L1GasPriceFactor = 0.04  
5  ByteGasCost = 16  
6  ZeroByteGasCost = 4  
7  NetProfit = 1.2  
8  BreakEvenFactor = 1.3  
9  FinalDeviationPct = 10  
10 L2GasPriceSuggesterFactor = 0.3
```

<https://github.com/0xPolygonHermes/zkevm-node/blob/develop/docs/config-file/node-config-doc.md>

- **L1GasPriceFactor**: is the percentage of the L1 gas price that will be used as the L2 min gas price
- **ByteGasCost**: cost per byte that is not 0.
- **ZeroByteGasCost**: cost per byte that is 0.
- **NetProfit**: is the profit margin to apply to the calculated breakEvenGasPrice.
- **BreakEvenFactor**: is the factor to apply to the calculated breakEvenGasPrice when comparing it with the gasPriceSigned of a tx.
- **FinalDeviationPct**: is the max allowed deviation percentage BreakEvenGasPrice on re-calculation
- **L2GasPriceSuggesterFactor**: is the factor to apply to L1 gas price to get the suggested L2 gas price used in the calculations when the effective gas price is disabled (testing/metrics purposes)