

# The zkEVM Architecture

## Part III: Communication Between Layers

---

Polygon zkEVM & Universitat Politècnica de Catalunya (UPC)

Marc Guzman-Albiol <marc.guzman.albiol@upc.edu>

Jose Luis Muñoz-Tapia <jose.luis.munoz@upc.edu>

Version 1.2

December 12, 2023

# Outline

Communication Between Layers

Exit Trees

An Efficient Append-only SMT

Building the Exit Trees for the zkEVM


The Global Exit Tree

Global Exit Tree Update in L1

Global Exit Tree Update in L2

Summary up to fork-dragonfruit (fork-5)

From fork-etrog (fork-6)




# The Bridge

## Bridge

The bridge is a subsystem of the zkEVM that is composed of several components and its main purpose is to enable **exchanges between different layers**.

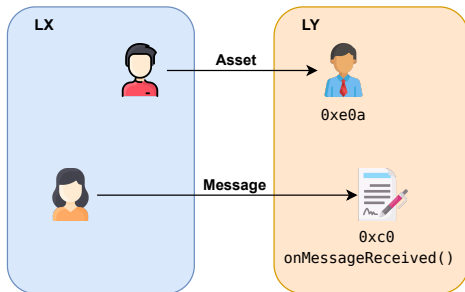
- We will start by defining exchanges between L1 and an L2 but our intention is to be general, that is, to enable exchanges between multiple layers LX and LY.
- This is why we call this subsystem the **LXLY bridge**.
- For the explanations, we will use three layers denoted as LX, LY and LZ.



# Exchanges: Assets and Messages


We enable the exchanges of:

- a) **Assets:** Ether or Tokens to accounts in the destination layer.
  - b) **Messages:** The execution of a function `onMessageReceived` of some contract.
- This is what we call the **messaging** mechanism of the bridge (messages can transfer Ether too).




# LXLY Bridge and Smart Contracts

- The core logic of the LXLY bridge is implemented in smart contracts.
- In particular, the main contract is a contract called **Bridge.sol** that is deployed in any layer in which we want exchanging enabled.
- One of our goals is that the **Bridge.sol** smart contract is **exactly the same** in all layers.




# LXLY Exchanges: Bridge-Claim Model i

The LXLY bridge follows a **bridge-claim** model:



As it can be observed, each bridge smart contract has the methods **bridge()** and **claim()**.

## LXLY Exchanges: Bridge-Claim Model ii



1. In the origin layer (e.g. LX), the user sends a transaction to the **bridge()** function providing the destination network (e.g. LY). Transactions to the bridge function are also known as "deposits".
2. In the destination layer (LY), the user sends a transaction to the **claim()** function providing the origin network (e.g. LX).

In the **Bridge.sol** smart contracts, we need a compact way of storing the information of calls to the bridge function (that we also call these data "exits" or "outgoing transmissions").



# Outline

Communication Between Layers

Exit Trees

An Efficient Append-only SMT

Building the Exit Trees for the zkEVM

The Global Exit Tree

Global Exit Tree Update in L1


Global Exit Tree Update in L2

Summary up to fork-dragonfruit (fork-5)

From fork-etrog (fork-6)

# Local Exit Tree


- In each layer, the bridge contract builds an **append-only Merkle tree** with all the exits, i.e. each call to the `bridge()` function.
- The Merkle tree with all the exits of a layer is called its **Local Exit Tree** and, its root is called **Local Exit Root (LER)**.
- Each leaf of a Local Exit Tree of a layer LX stores a single exit.



**Remark.** An exit tree is a different object and has a different structure than the tree that stores the L2 state.

# Leaves of Exit Trees i

Local Exit Tree



Each leaf contains:

```
1 uint8 leafType,  
2 uint32 originNetwork,  
3 address originAddress,  
4 uint32 destinationNetwork,  
5 address destinationAddress,  
6 uint256 amount,  
7 bytes32 metadataHash
```

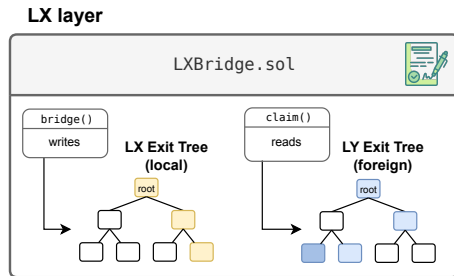
## Leaves of Exit Trees ii

- **leafType** used to identify whether the leaf is an asset or a message:
  - Asset: value is 0.
  - Message: value is 1.
- **originNetwork**: the identifier (**chainId**) of origin layer of the exchange.
- **originAddress**: if it is an asset exchange, it is the address of the token contract.  
If it is an message exchange, it is the source address of the bridge call.
- **destinationNetwork**: the identifier of the destination layer (**chainId**) of the exchange.
- **destinationAddress**: is the account receiving the asset or the address of the smart contract if it is a message exchange.
- **leafAmount**: amount of asset exchanged (Ether or Tokens).
- **bytes32 metadataHash**: the hash of the metadata.
  - Asset: the metadata is the name, symbol and decimals of the token.
  - Message: the metadata is the **calldata** for calling the **onMessageReceived()** function.

# Read and Write Exit Trees

In each layer, the corresponding bridge smart contract needs to:

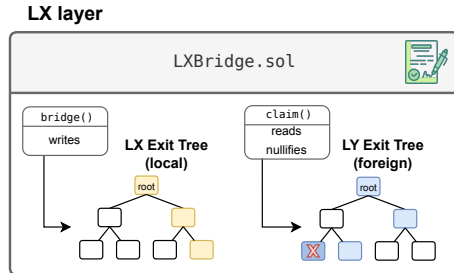
- **Write:** write in its Local Exit Tree each new exit resulting from a call to its `bridge()` function.
- **Read:** read the Exit Trees of other layers to process calls to `claim()`.



# Read and Nullify

In each layer, the corresponding bridge smart contract needs to additionally **nullify** the claim:

- Each claim must be **locally nullified** in the bridge contract to avoid double claimings.
- The nullify process to avoid claiming transactions that have already been processed uses an efficient mapping known as **claimedBitMap** (we will explain how this bitmap works later in more detail).



# Outline

Communication Between Layers

Exit Trees

**An Efficient Append-only SMT**

Building the Exit Trees for the zkEVM

The Global Exit Tree

Global Exit Tree Update in L1

Global Exit Tree Update in L2

Summary up to fork-dragonfruit (fork-5)

From fork-etrog (fork-6)

# An Efficient Append-only SMT

The idea is to create an append-only sparse tree whose successive roots can be computed with a minimal amount of persistent data.

- It turns out that to append new data elements to the tree we are going to just need to store:
  1. An array of the size of the tree depth (denoted as `branch`).
  2. The last appended element's index (denoted as `lastElemIndex`).
- The depth of the tree (maximum capacity) will be known a priori.
- As a result, we will not need to use markers for branches and leaves.
- Furthermore, the tree will be balanced.



# Empty Append-only SMT i

For our append-only tree we are going to use 0s as default value for empty leaves.  
When the incremental Merkle tree is empty, we have that:


$$h_i = 0$$

$$h_{j,k} = h(0|0) = h^{(00)}$$


$$h_{m,n,\ell} = h(h^{(00)}|h^{(00)}) = h^{(0000)}$$

...

- So, we just need to compute a different hash value per level.
- Let's consider as a toy example a small incremental Merkle tree of a maximum capacity of 8 leaves ( $d = 3$ ).



# Empty Append-only SMT ii



`lastElemIndex = 0`   `branch = [0, 0, 0]`

$$\text{root} = h^{(00000000)}$$

$$h^{(00000000)} = h(h^{(0000)} | h^{(0000)})$$


$$h^{(0000)} = h(h^{(00)} | h^{(00)})$$

$$h^{(00)} = h(0 | 0)$$

- Note that we can compute the root of the empty tree from the zero hash values.
- Note also that the hash of zero nodes is uniquely determined by the height of the subtree under the node.

# Append the First Element

Let's suppose that we want to add an element  $x_0$ :



`lastElemIndex = 1`   `branch = [ $h_0$ , 0, 0]`

`root =  $h_{01234567}$`

$h_{01234567} = h(h_{0123} | h^{(0000)})$

$h_{0123} = h(h_{01} | h^{(00)})$

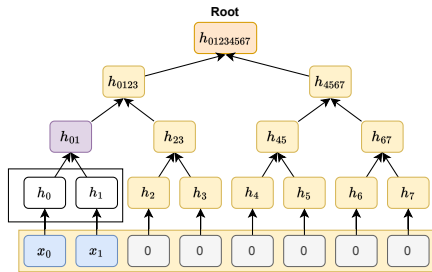
$h_{01} = h(h_0 | 0)$

$h_0 = h(x_0)$

- Note that we "can" (lazy) compute the root of the current tree with the `branch` and `lastElemIndex`.
- Just need to write `branch[0] =  $h_0 = h(x_0)$` .

# Append an Second Element

Let us now add another element  $x_1$  into the tree:



`lastElemIndex = 2`   `branch = [h0, h01, 0]`

`root = h01234567`

$h_{01234567} = h(h_{0123} | h^{(0000)})$

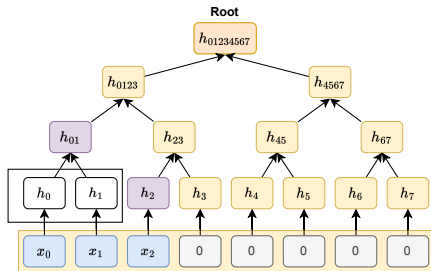
$h_{0123} = h(h_{01} | h^{(00)})$

$h_{01} = h(h_0 | h(x_1))$

- Note that we can compute the root of the current tree by writing  
`branch[1] = h01 = h(h0 | h(x1))` in the **branch**.
- From now on,  $h_0$  and  $h_1$  are not needed any more for updating the root since they are integrated in  $h_{01}$ .

# Append a Third Element

Let us now add another element  $x_2$  into the tree:



`lastElemIndex = 3`   `branch = [ $h_2$ ,  $h_{01}$ , 0]`

`root =  $h_{01234567}$`

`$h_{01234567} = h(h_{0123} | h^{(0000)})$`

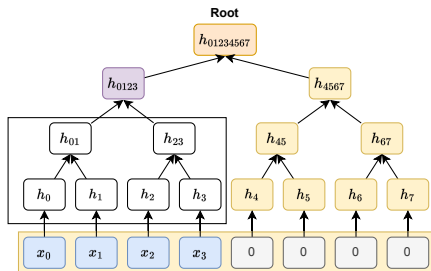
`$h_{0123} = h(h_{01} | h_{23})$`

`$h_{23} = h(h_2 | 0)$`

`$h_2 = h(x_2)$`

# Append a Fourth Element

Let us now add another element  $x_3$  into the tree:



`lastElemIndex = 4`   `branch = [h2, h01, h0123]`

`root = h01234567`

$h_{01234567} = h(h_{0123} | h^{(0000)})$

$h_{0123} = h(h_{01} | h_{23})$

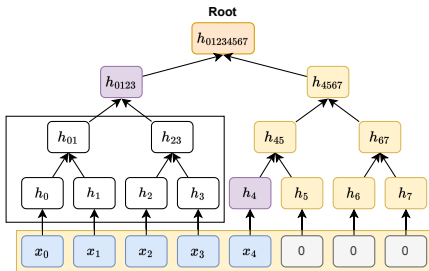
$h_{23} = h(h_2 | h_3)$

$h_3 = h(x_3)$

From now on,  $h_0, h_1, h_2, h_3, h_{01}$  and  $h_{23}$  are not needed any more for updating the root since they are integrated in  $h_{0123}$ .

# Append a Fifth Element

Let us now add another element  $x_4$  into the tree:



`lastElemIndex = 5`   `branch = [ $h_4$ ,  $h_{01}$ ,  $h_{0123}$ ]`

`root =  $h_{01234567}$`

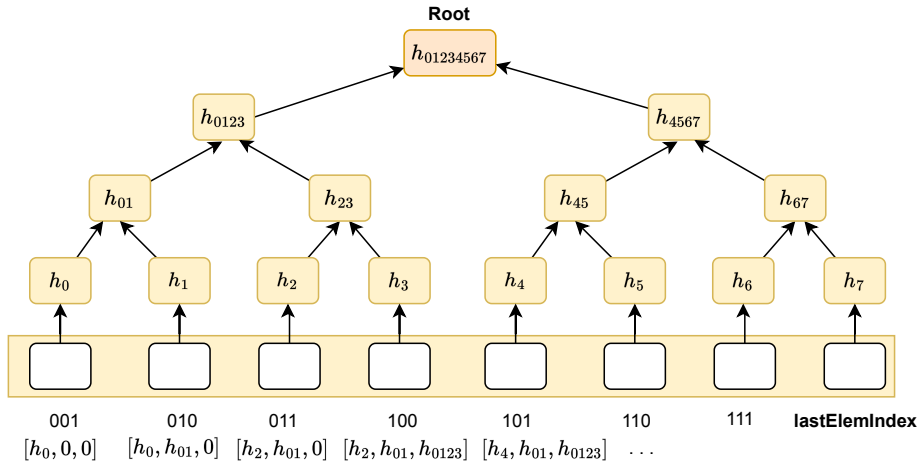
$h_{01234567} = h(h_{0123} | h_{4567})$

$h_{4567} = h(h_{45} | h^{(00)})$

$h_{45} = h(h_4 | 0)$

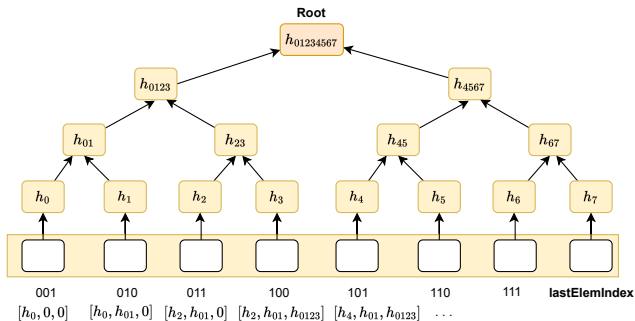
$h_4 = h(x_4)$

# The Written Position in the **branch i**






## The Written Position in the **branch** ii

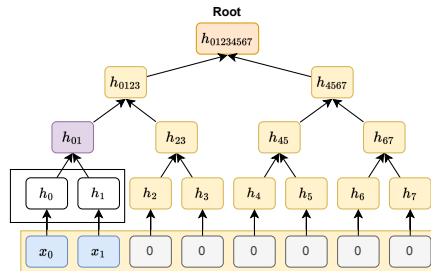


In general, the position of the **branch** that we have to update corresponds to the position of the **less significant bit to 1** of the binary representation of **lastElemIndex**.

# The Written Position in the **branch** iii

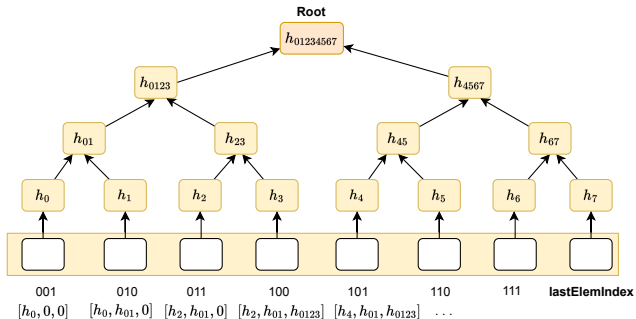


lastElemIndex = 1 (0b001) branch = [ $h_0$ , 0, 0]



lastElemIndex = 2 (0b010) branch = [ $h_0$ ,  $h_{01}$ , 0]

# Writing the branch i




The following pseudo-code computes the value and position to write in the **branch** array:

```
currentHash = dataHash
index = 0

while read bit of lastElemIndex next lsb {
  if(bit == 0) {
    currentHash = h(branch[index]|currentHash)
  } else { // bit == 1
    branch[index]= currentHash
    return
  }
  index++
}
```

## Writing the branch ii




lastElemIndex = 1 (0b001) branch = [ $h_0$ , 0, 0]

```
currentHash = dataHash
index = 0

while read bit of lastElemIndex next lsb {
  if(bit == 0) {
    currentHash = h(branch[index]|currentHash)
  } else { // bit == 1
    branch[index]= currentHash
    return
  }
  index++
}
```

## Writing the branch iii




lastElemIndex = 2 (0b010)    branch =  $[h_0, h_{01}, 0]$

```
currentHash = dataHash
index = 0

while read bit of lastElemIndex next lsb {
  if(bit == 0) {
    currentHash = h(branch[index]|currentHash)
  } else { // bit == 1
    branch[index]= currentHash
    return
  }
  index++
}
```

# Computing the Root i



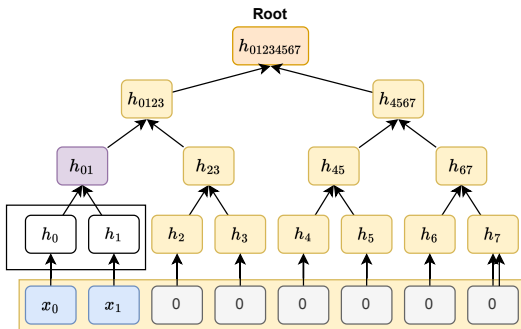
lastElemIndex = 1 (0b001)  
branch = [ $h_0$ , 0, 0]

index = 2 bit = 0  
currentHash =  $h(\text{currentHash} \mid \text{currentZeroHash}) = \text{root}$

index = 1 bit = 0  
currentHash =  $h(\text{currentHash} \mid \text{currentZeroHash}) = h_{0123}$   
currentZeroHash =  $h(\text{currentZeroHash} \mid \text{currentZeroHash}) = h^{(0000)}$

index = 0 bit = 1  
currentHash =  $h(\text{branch}[\text{index}] \mid \text{currentHash}) = h_{01}$   
currentZeroHash =  $h(\text{currentZeroHash} \mid \text{currentZeroHash}) = h^{(00)}$   
currentZeroHash = 0  
currentHash = currentZeroHash

# Computing the Root ii



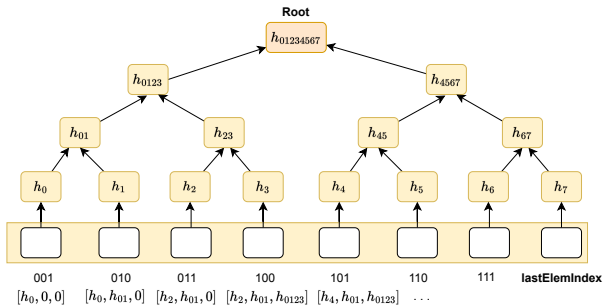
index = 2 bit = 0  
currentHash =  $h(\text{currentHash} \mid \text{currentZeroHash}) = \text{root}$

index = 1 bit = 1  
currentHash =  $h(\text{branch}[\text{index}] \mid \text{currentHash}) = h_{0123}$   
currentZeroHash =  $h(\text{currentZeroHash} \mid \text{currentZeroHash}) = h^{(0000)}$

index = 0 bit = 0  
currentHash =  $h(\text{currentHash} \mid \text{currentZeroHash}) = h_{23}$   
currentZeroHash =  $h(\text{currentZeroHash} \mid \text{currentZeroHash}) = h^{(00)}$

currentZeroHash = 0  
currentHash = currentZeroHash

# Computing the Root iii



The following pseudo-code computes the root:

```
currentZeroHash = 0
currentHash = currentZeroHash
index = 0

while read bit of lastElemIndex next lsb {
  if(bit == 0) {
    currentHash = h(currentHash | currentZeroHash)
  } else { // bit == 1
    currentHash = h(branch[index] | currentHash)
  }
  currentZeroHash = h(currentZeroHash | currentZeroHash)
  index++
}
```



## Some Properties of the Append-only SMT

- Note that we built a key-value tree that is sparse, binary and balanced.
- We also applied partial tree construction and we did lazy evaluation in the computation of the root.
- Insertions only need to write one position in the **branch** array.
- Finally, we could “revoke” or “nullify” elements if we use an extra structure to denote which elements of the append-only tree are no longer valid.

# Outline

Communication Between Layers

Exit Trees

An Efficient Append-only SMT

Building the Exit Trees for the zkEVM

The Global Exit Tree

Global Exit Tree Update in L1

Global Exit Tree Update in L2

Summary up to fork-dragonfruit (fork-5)

From fork-etrog (fork-6)

# Smart Contract for the Exit Tree

The smart contract to manage the exits can be found at:

<https://github.com/0xPolygonHermes/zkevm-contracts/blob/main/contracts/lib/DepositContract.sol>

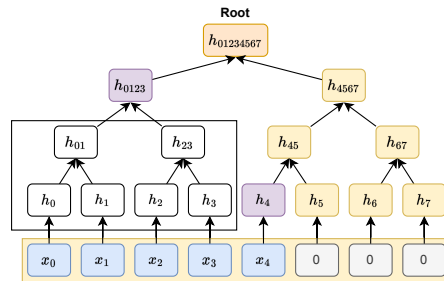
The contract is based on the implementation of the Eth 2.0 deposit contract.

The contract has four functions:

- `function getLeafValue(...)` which computes the hash of a leaf.
- `_deposit(bytes32 leafHash)` which adds a new leaf to the tree
- `getDepositRoot()` which computes the Merkle root.
- `function verifyMerkleProof(...)` which verifies a Merkle proof.

# Obtaining Merkle Proofs of Previous Deposits

- Notice that since the branch is updated, the smart contract is not able to provide the Merkle proof for old deposits.
- In the example, the Merkle proof for  $x_0$  is the tuple  $(h_1, h_{23}, h_{4567})$  which is not currently stored on the smart contract.



$lastElem = 5$     $branch = [h_4, h_{01}, h_{0123}]$

- If we want to simplify the process of obtaining a Merkle proof for users, we have to store all the Merkle tree nodes and provide a service that answers with the appropriate Merkle proofs.
- This service in the zkEVM is called the **claim service**.
- The implementation is called claim service manager and it is available at:

<https://github.com/0xPolygonHermes/zkevm-bridge-service/tree/develop/claimtxman>

# Outline

Communication Between Layers

Exit Trees

An Efficient Append-only SMT

Building the Exit Trees for the zkEVM

**The Global Exit Tree**

Global Exit Tree Update in L1

Global Exit Tree Update in L2

Summary up to fork-dragonfruit (fork-5)


From fork-etrog (fork-6)

# The Global Exit Tree

So, while processing a `claim()` at a certain layer, this layer might need to read an exit leaf stored in the Exit Tree of another layer.


For enabling this, we build a Merkle tree that includes all the exit trees of all the layers.

- In particular, the local exit roots are used as leaves to build the **Global Exit Tree**.
- The root of the Global Exit Tree will be called **globalExitRoot** and it is a cryptographic digest of all the Local Exit Roots it contains.
- With the **globalExitRoot** and the appropriate **Merkle proofs**, a bridge contract can securely **read a leaf from any exit tree**.




# Updating the `globalExitRoot`

- From the point of view of a layer, the Global Exit Tree can change because:
  - The local exit tree of the layer changes.
  - The local exit tree of another layer changes.



LX: Local Exit Root Update



LX: Foreign Exit Root Update



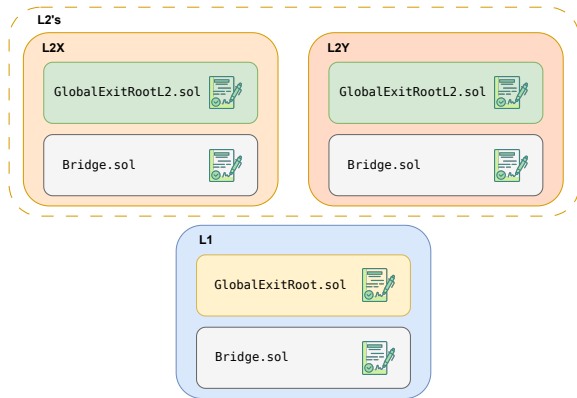
## The `globalExitRoot` Update is Layer-specific i

- Each Bridge smart contract does the update of its local exit root with the same logic when there are calls to the `bridge()` function.
- However, as we will see next, the `globalExitRoot` update is performed differently in the bottom layer (L1) than in a top layer (L2) when there is an update of a foreign exit root.
- In other words, the logic to update the global exit root is layer-specific.
- But, recall that we want to deploy the same bridge smart contract in all the layers.
- So, to achieve this, the solution is to deploy the logic for updating the `globalExitRoot` in to a separate smart contract in each layer.

# The `globalExitRoot` Update is Layer-specific ii

The contract that manages the update of the `globalExitRoot` in the bottom layer (L1) is called `GlobalExitRoot.sol`.

The contract that manages the update of the `globalExitRoot` in a top layer is called `GlobalExitRootL2.sol`.



# Outline

Communication Between Layers

Exit Trees

An Efficient Append-only SMT

Building the Exit Trees for the zkEVM

The Global Exit Tree

Global Exit Tree Update in L1

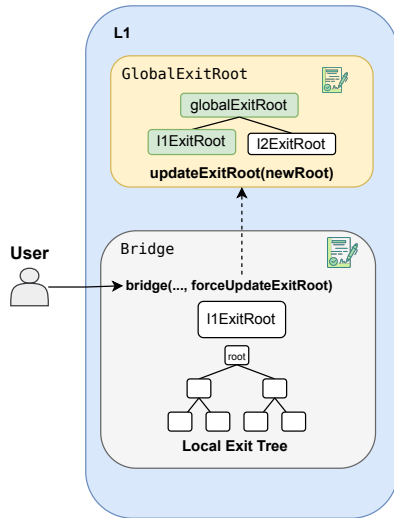
Global Exit Tree Update in L2

Summary up to fork-dragonfruit (fork-5)

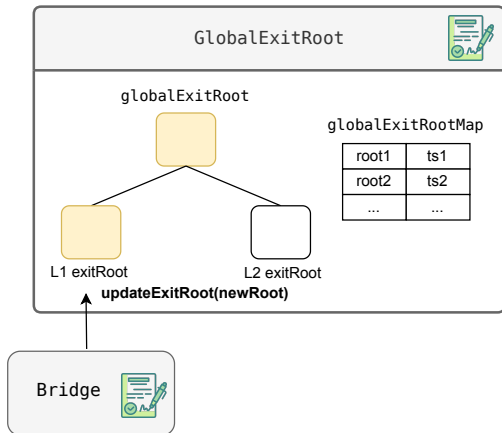
From fork-etrog (fork-6)

# Updating the `globalExitRoot` in L1: Local Update

- When a user sends a transaction with a call to `bridge()` (i.e. if someone makes a deposit) the corresponding local exit root (in this case L1) is modified.
- The local exit root is managed by the Bridge smart contract.
- The execution of the bridge (deposit) **may update** the `globalExitRoot`, which is stored in the `GlobalExitRoot` smart contract, by calling the `updateExitRoot()` function.
- If the bool parameter `forceUpdateExitRoot` is set as true when calling the `bridge()` function, the bridge smart contract calls `updateExitRoot()` in the `GlobalExitRoot` contract to perform the update.



`GlobalExitRoot.sol` contains a mapping that associates a `globalExitRoot` with the update timestamp, providing a mechanism to keep track of each root update.



## About the `forceUpdateExitRoot` flag


- The bool `forceUpdateExitRoot` can be set as:
  - `true`: the transaction costs around 120k gas.
  - `false`: the transaction costs around 40k gas (about 30% of the cost with the flag activated).
- The idea of the `updateGlobalExitRoot()` function in `Bridge.sol` is to be able to generate a sponsored `globalExitRoot` update service.

## More Ways of Doing the `globalExitRoot` Local Update

- Let's consider that happens if no user calls the `bridge()` function with `forceUpdateExitRoot` flag equal to `true`.
- Then, the question is who updates the `globalExitRoot` in the L1 `GlobalExitRoot` smart contract and when.
- In the zkEVM architecture, two actors can perform that job:
  - a) **Anyone** can send a transaction to the `updateExitRoot()` function of the L1 `Bridge` smart contract.
  - b) The **sequencer**, when calling `sequenceBatches()` function of the `ZkEVM` smart contract at the end calls the `updateExitRoot()` function of the `Bridge` contract.

## Updating the `globalExitRoot` in L1: Remote Update


- Suppose there is an outgoing transmission from Layer 2, that is, an L2 transaction calling the `bridge()` function of the **Bridge** smart contract in L2.
- This L2 transaction changes the state of the L2 exit tree.
- Q: How does L1 realizes that the L2 `exitRoot` has changed?





# Re-engineering the Proof i

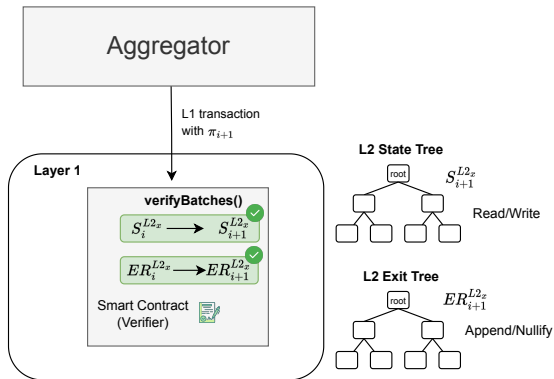
To make the state of the L2 **exitRoot** available to L1, we add a new parameter in the proof that contains the new L2 **exitRoot** that results after processing the batch.



# Re-engineering the Proof ii

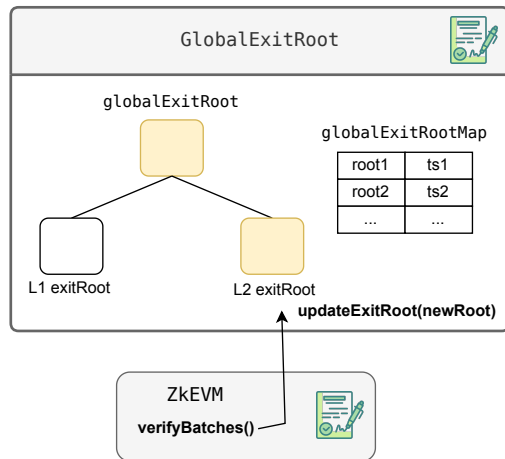
So, our new proof now assures that:

1. The processing of the L2 transactions is correct and so it is the new L2 state root (`newStateRoot`).
2. The new L2 `exitRoot` (`newLocalExitRoot`) is also correct.



# Processing the Proofs in L1

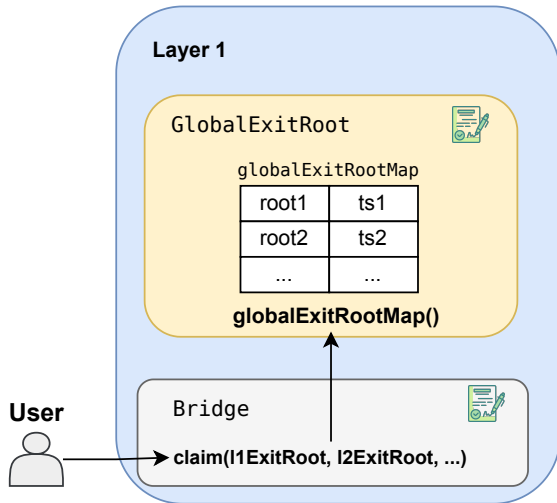
- The `verifyBatches()` function of `ZkEVM.sol` is who verifies the proofs.
- If the proof is correct, at the end will call the `updateExitRoot()` of the `GlobalExitRoot.sol` smart contract to perform the `globalExitRoot` update (providing the exit three of the foreign layer).
- As previously mentioned, all the `globalExitRoots` are registered in a mapping in the `GlobalExitRoot.sol` smart contract.
- We can obtain the latest root calling `getLastGlobalExitRoot()` at `GlobalExitRoot.sol`.



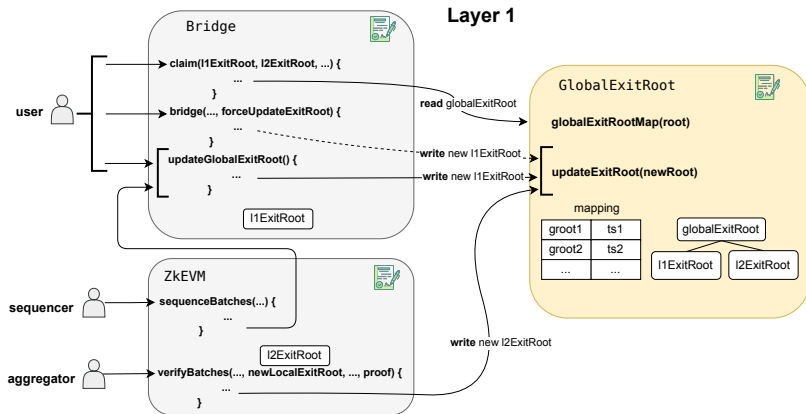
<https://github.com/0xPolygonHermez/zkevm-contracts/blob/main/contracts/PolygonZkEVMGlobalExitRoot.sol>

# Reading the `globalExitRoot` in L1

When a user sends a transaction with a call to `claim()`, that is to say, when someone makes a withdraw, the **Bridge** smart contract needs to read a valid `globalExitRoot` to perform the claim.



# Summary of the **globalExitRoot** Management in L1



# Outline

Communication Between Layers

Exit Trees

An Efficient Append-only SMT

Building the Exit Trees for the zkEVM

The Global Exit Tree

Global Exit Tree Update in L1

Global Exit Tree Update in L2

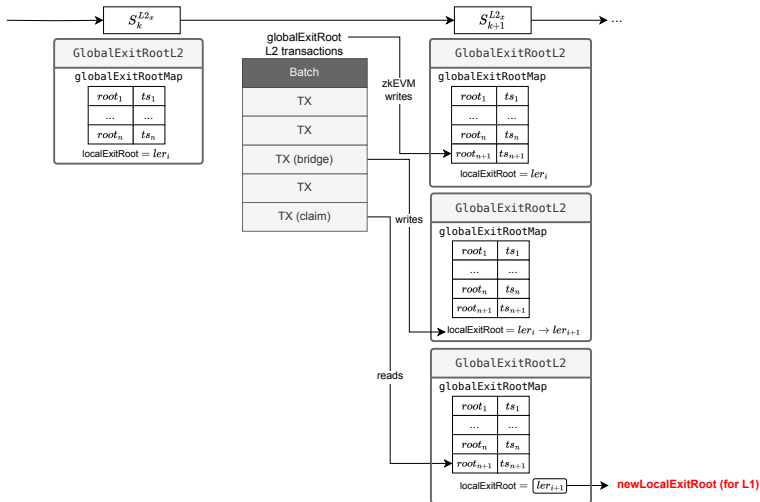
Summary up to fork-dragonfruit (fork-5)

From fork-etrog (fork-6)

# Overview of Exit Trees Updates in L2 i

- The L2 also needs to be aware of changes in exit trees, that is, we have to manage local and remote updates of the exit trees in L2.
- In summary, this is done as follows:
  - The L2 local `exitRoot` is updated for each successful deposit (L2 transaction to `bridge()`).
  - The `globalExitRoot` is updated only **at the beginning of the processing of the L2 batch**.
  - Each batch, apart from the L2 transactions, includes a `globalExitRoot`.
  - The zkEVM processing of the batch inserts this `globalExitRoot` in the `globalExitRootMap` of the `GlobalExitRootL2` contract instance at L2.
  - The new `globalExitRoot` or any other previous root available in the `globalExitRootMap` can be used for doing the L2 claim transactions.

# Overview of Exit Trees Updates in L2 ii

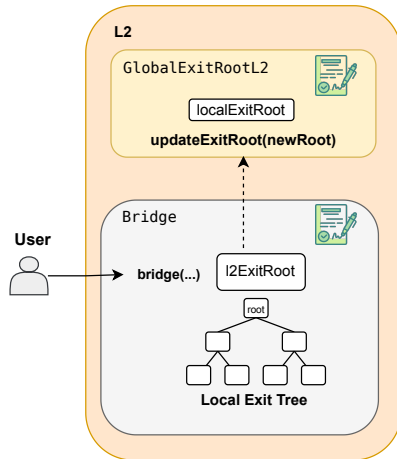




# Performing bridges (deposits) in L2

Processing transactions that call **bridge()** in L2 is identical to what we do in L1:

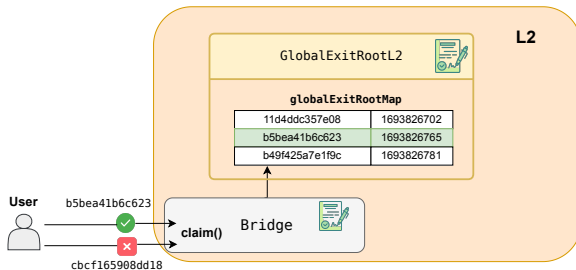
- When a user sends an L2 transaction with a call to **bridge()**, that is, when someone makes a deposit, the corresponding local exit root (in this case L2) is modified.
- If the bool parameter **forceUpdateExitRoot** is set as true when calling the **bridge()** function, the **Bridge** smart contract calls **updateExitRoot()** in the **GlobalExitRootL2** contract to perform the update.



# Performing claims (withdraws) in L2

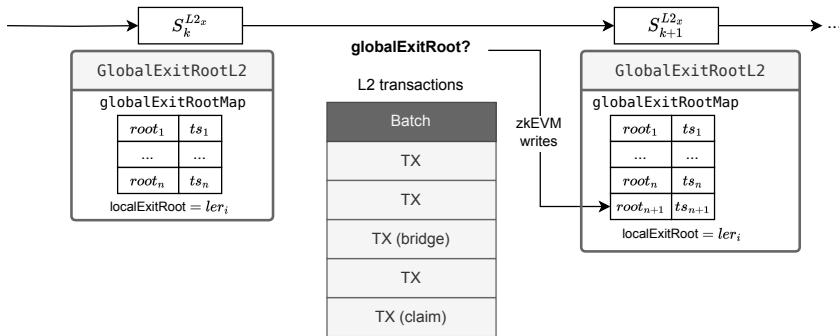
To do a claim in L2, we need to access a `globalExitRoot` that includes the data of our deposit in the origin network (L1).

- To make updated `globalExitRoots` available to L2, a `globalExitRoot` is provided with each batch.
- At the beginning of the ROM processing of a L2 batch, the executor inserts the `globalExitRoot` (and the corresponding timestamp) received as an input in the `globalExitRootMap` of `GlobalExitRootL2` contract.
- Finally, the **Bridge** smart contract in L2 can check if the deposit exists if the user provides a valid Merkle proof against a global exit root that is registered in the mapping of the `GlobalExitRootL2` contract.



# Who Decides the **globalExitRoot** for L2? i

Who decides the **globalExitRoot** that goes with the batch?



## Who Decides the `globalExitRoot` for L2? ii

- Note that if a `globalExitRoot` is registered or not influences the result of the batch execution.
- In more detail, depending on the available `globalExitRoots`, some claims will be valid and some won't.
- Recall that sequencer does a batch pre-execution to check if the batch fits in the available resources so, the sequencer must deterministically know the result of the execution.
- Therefore, the sequencer is the actor that must associate a `globalExitRoot` for each batch.
- When sequencing the batch, the `ZkEVM` contract checks that the `globalExitRoot` provided by the sequencer exists in the mapping of the L1 `GlobalExitRoot` contract.
- Notice that the sequencer decides the `globalExitRoot` but, the aggregator also needs to read this value to generate the corresponding proof.

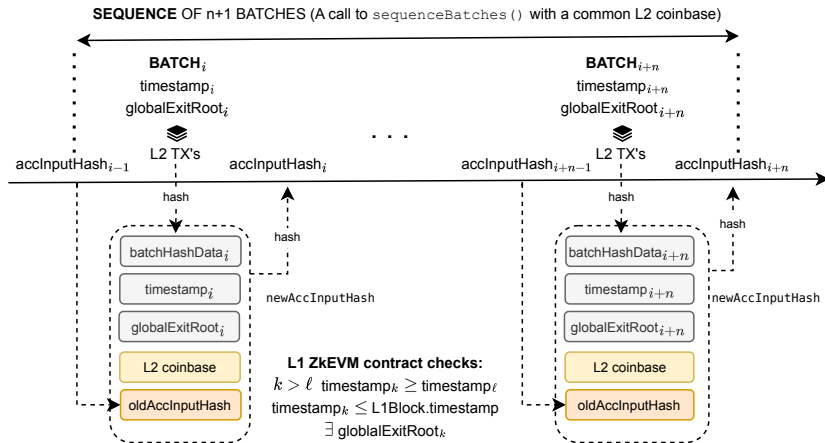
## How is the **globalExitRoot** for the Batch Available to the Aggregator?

- The aggregator needs to use the same **globalExitRoot** decided by the sequencer to build the proof of a batch.
- To achieve this, the **globalExitRoot** is passed by the sequencer as part of the batch data when the batch is sequenced:

```
1 struct BatchData {  
2     bytes transactions;  
3     bytes32 globalExitRoot;  
4     uint64 timestamp;  
5 }
```

- Then, the smart contract includes this parameter as part of the hashed input data to build the cryptographic pointer for the batch (**accInputHash**).

# Sequences of Batches i



- The zkEVM smart contract checks that:
  - a) The **timestamps** of batches are correct:
    - The timestamps of a batch is set by the sequencer when the batch is started being filled with transactions.
    - Then, for batches in the sequence with  $k > \ell$ , then:

$$timestamp_k > timestamp_\ell.$$

- All timestamps in the sequence are smaller than the timestamp of the L1 sequencing transaction:

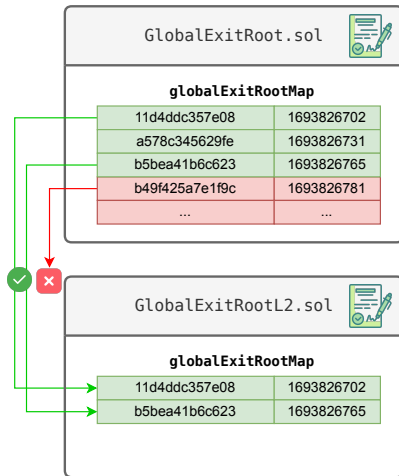
$$timestamp_k < block.timestamp.$$

- Recall that in fork-dragonfruit (fork-5) all the blocks (transactions) within a block share the same timestamp.

- b) The **globalExitRoot** provided for each batch exists.

# Censoring the **globalExitRoot** Progress (Avoids Withdraws)

- Notice that a sequencer may refuse to make the **globalExitRoot** progress.
- The situation is that there are new **globalExitRoots** available in the **GlobalExitRoot** contract at L1, but the sequencer refuses to use them for sequenced batches.
- Notice that a user can do a claim using any already registered global root.
- But the withdraw of a deposit of a user can be effectively censored if the sequencer, when sequencing the batches, does not include an updated **globalExitRoot** that includes the data of the deposit.





# Countermeasure for the Censoring of the `globalExitRoot`

- To fix the problem of the censorship of the `globalExitRoot` we can use the same mechanism as we use to avoid the censorship of L2 transactions.
- This mechanism is the so called "forced batches".
- In a forced batch you can include the transactions that you like and also the `globalExitRoot` that you like:

```
1 struct ForcedBatchData {  
2     bytes transactions;  
3     bytes32 globalExitRoot;  
4     uint64 minForcedTimestamp;  
5 }
```

- In this way, you can enable your deposits for their corresponding withdraws.

# Outline

Communication Between Layers

Exit Trees

An Efficient Append-only SMT

Building the Exit Trees for the zkEVM

The Global Exit Tree

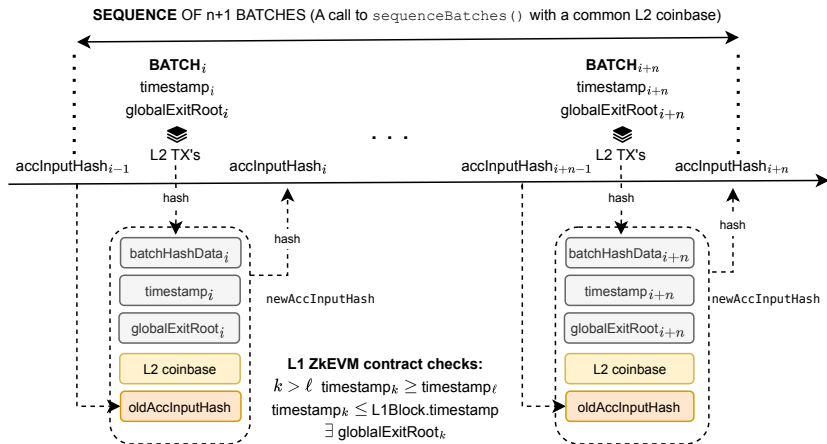
Global Exit Tree Update in L1

Global Exit Tree Update in L2

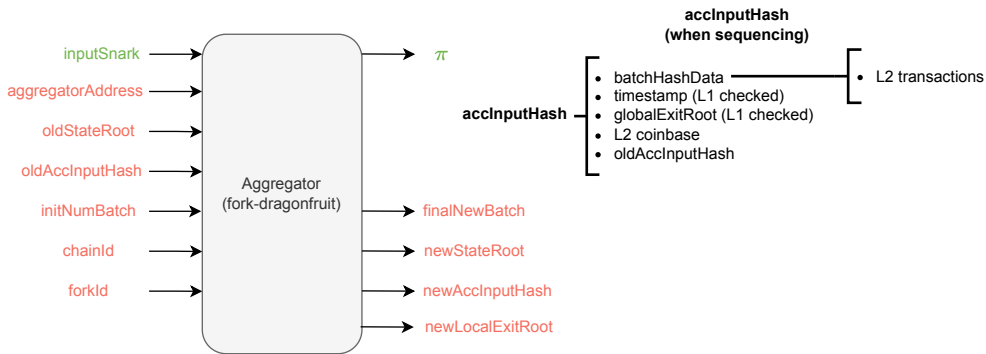
Summary up to fork-dragonfruit (fork-5)

From fork-etrog (fork-6)

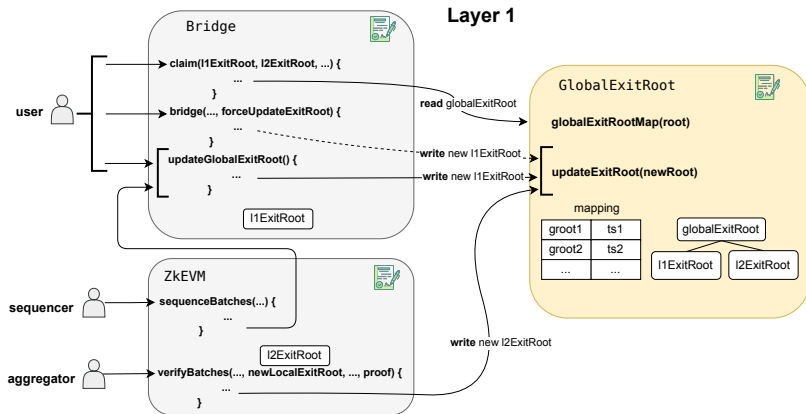
# Sequencing




# Proving and Sequencing



# globalExitRoot in L1



# globalExitRoot in L2



# Outline

Communication Between Layers

Exit Trees

An Efficient Append-only SMT

Building the Exit Trees for the zkEVM

The Global Exit Tree

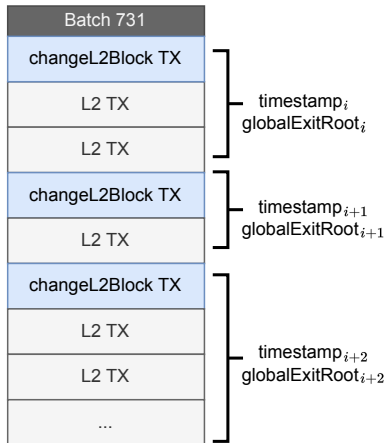
Global Exit Tree Update in L1

Global Exit Tree Update in L2

Summary up to fork-dragonfruit (fork-5)

From fork-etrog (fork-6)

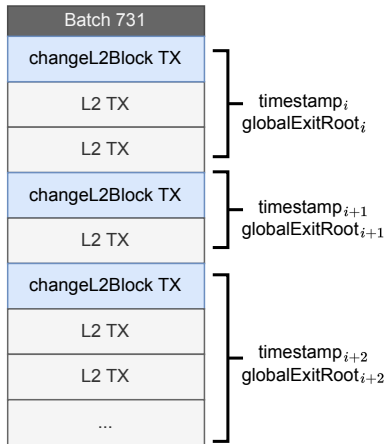
## L2 Blocks and the changeL2Block Transaction



- In the fork-dragonfruit, the blocks within a batch share the same timestamp and `globalExitRoot`.
- In fork-etrog, each block has:
  - Its own `timestamp`.
  - Its own `globalExitRoot`.
- These parameters are provided by the sequencer in the "special" transaction `changeL2Block`.
- Each batch must start with a `changeL2Block` transaction.



# Moving Checks from L1 to zkEVM Processing



- In fork-dragonfruit, the checks over the Batch's **timestamp bounds** and the **globalExitRoot existence** were performed by the L1 ZkEVM smart contract.
- In fork-etrog, notice that since we have a different timestamp (and possibly **globalExitRoot**) per L2 block, we have much more checks to perform.
- To decrease L1 costs, we move the checks of the **globalExitRoot existence** and the **timestamp bounds** to the zkEVM processing.
- That is, these checks are included in the proof and removed from the ZkEVM L1 smart contract.

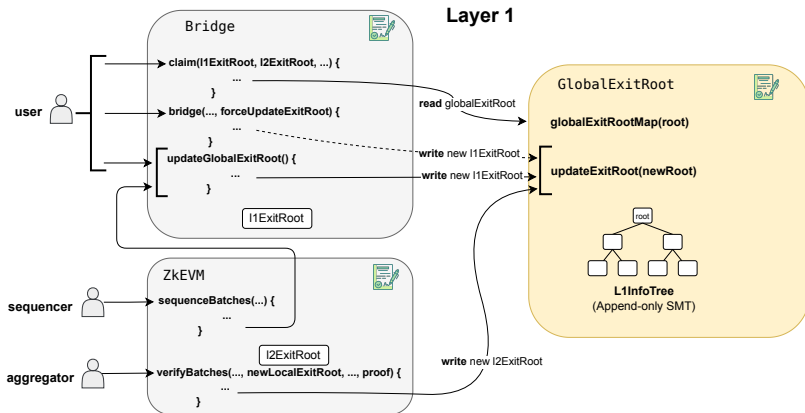
# Checking the `globalExitRoot` Existence at the zkEVM Processing

## `globalExitRoot` existence

To check the existence of a `globalExitRoot`, the zkEVM proving system needs to **have access to all the `globalExitRoots` recorded at L1.**

- However:
  - The `globalExitRoots` are stored in a mapping in L1, at the `GlobalExitRootManager` smart contract, however, we cannot pass a mapping to the prover.
  - We could think in passing the list of `globalExitRoots` but this is inefficient since this list is a potentially big and always growing data structure.
- The best way is to build a Merkle tree with all the `globalExitRoots`.
- This tree is called the **L1InfoTree**.


# globalExitRoot in L1: L1InfoTree



The **L1InfoTree** is an append-only SMT with same implementation as exit trees that is updated with new `globalExitRoots` by L1 `GlobalExitRoot` contract (replaces mapping of fork-dragonfruit).

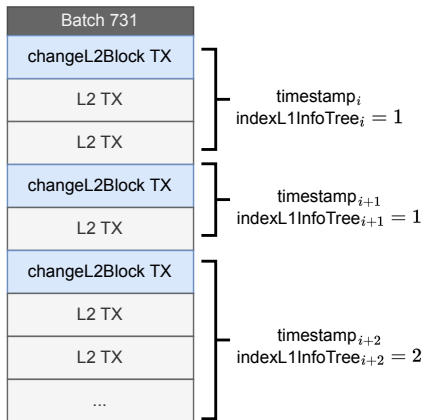
# Proving Batches using the L1InfoTree $i$

- The prover needs to have access to:
  1. The root of the L1InfoTree.
  2. The index of the **globalExitTree** being used for processing the L2 block.
- The index of the **globalExitTree** being used is called **indexL1InfoTree** and it is provided in the **changeL2Block** transaction.

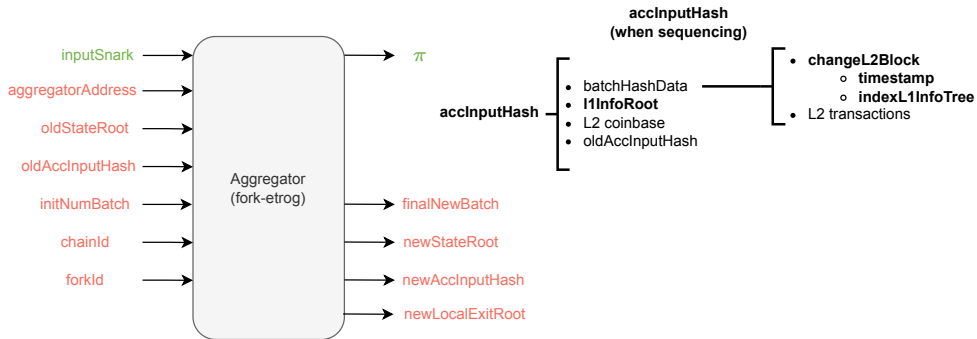


indexL1InfoTree = 1   indexL1InfoTree = 2

**L1InfoTree**  
(Append-only SMT)



# Proving Batches using the L1InfoTree ii

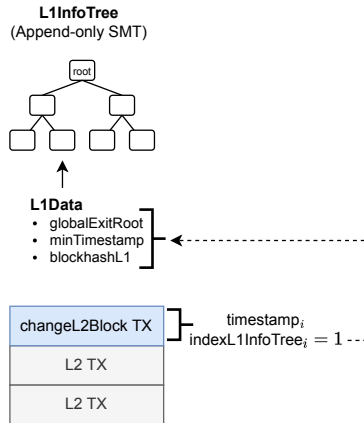


Note. Aggregator needs to know the SMT proofs given (indexL1InfoTree, l1InfoRoot) to compute the proof.

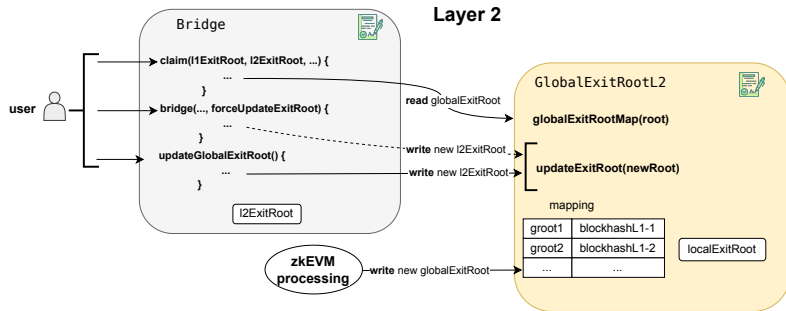
# L1Data and the blockhashL1

In fact, each leaf of the **L1InfoTree** stores the following information:

- **globalExitRoot**.
- **minTimestamp**:
  - It is the time at which the **globalExitRoot** was recorded.
  - It is used by the timestamp checks as the minimum timestamp possible for a block (explained next).
- **blockhashL1**:
  - Blockhash of the L1 block that precedes the block in which it is placed the transaction that inserts the **globalExitRoot** in the **L1InfoTree**.
  - Recall that the header of an Ethereum block includes the (L1) state root, so making available the **blockhashL1** provides the L1 state to L2 contracts.

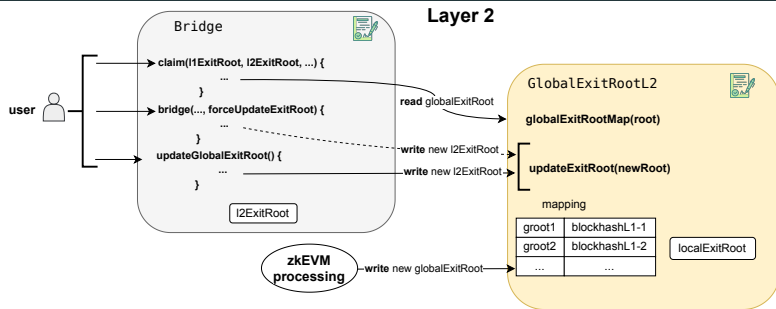


# globalExitRoot in L2 i



- The zkEVM processing inserts the new `globalExitRoots` in the mapping of the `GlobalExitRootL2` contract, however, the mapping is **not updated** if:
  - The `globalExitRoot` is already inserted or,
  - The `indexL1InfoTree` is 0 (the first element of the tree does not contain data but its index has the special purpose of not upgrading and saves gas and data availability).

# globalExitRoot in L2 ii



- Notice that unlike in fork-dragonfruit in which the mapping stored timestamps, in fork-etrog we store the **blockhashL1** associated with the **globalExitRoot**.
- The **blockhashL1** is also stored as part of the **blockhashL2**, which provides a summary of the execution of the L2 block including the current L2 state.
- The **blockhashL1** can be used by L2 transactions, during their processing, to access L1 data.



# The `changeL2Block` Transaction i

Field Name	Size
<code>type</code>	1 byte
<code>deltaTimestamp</code>	4 bytes
<code>indexL1InfoTree</code>	4 bytes

The `type` field:

- It is used to distinguish the `changeL2Block` transaction from regular L2 transactions.
- The value used is `0x0C`, while regular L2 transactions are rlp-encoded and their first byte is always bigger than `0xC0`.
- We also leave room for Ethereum typed transactions which use low values in their type field.

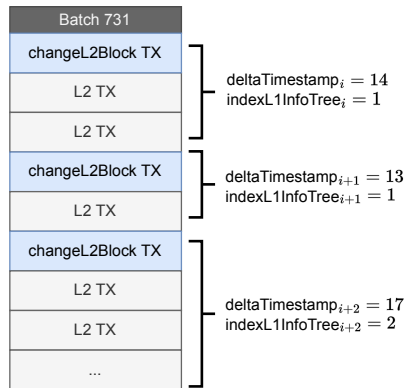
The `indexL1InfoTree` field:

- This is the index of the `globalExitRoot` being used by the block.
- The `L1InfoTree` has 32 levels, that is, keys of 32 bits (4 bytes).
- Recall that 0 has the special meaning of not updating in L2.

# The `changeL2Block` Transaction ii

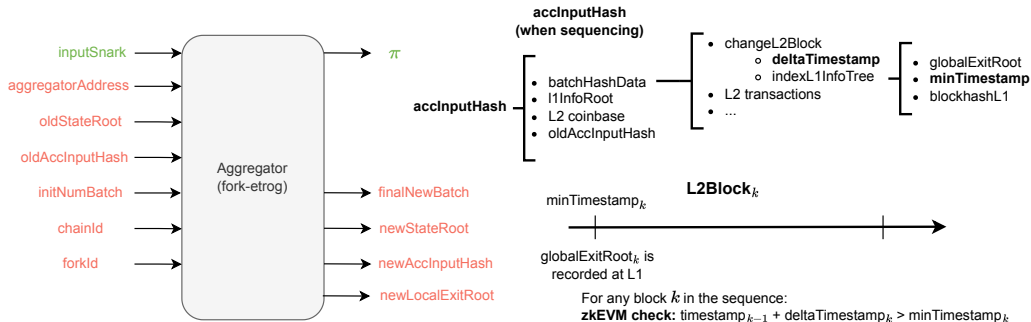
The `deltaTimestamp` field:

- Instead of using absolute timestamps, we use incremental timestamps.
- The `deltaTimestamp` shows the amount of seconds that need to be **added to the timestamp of the previous L2 block** to obtain the timestamp of the current block.
- The `timestamp` of the previous L2 block is available to the zkEVM in the system contract `0x5ca1ab1e` as part of the `blockhashL2`.



**Note.** We use incremental timestamps to reduce the size of this field (data availability): using a regular Unix time timestamp we would need to use 64 bits, while increments are always much smaller and we use 32-bits.

# Checking the Lower timestamp Bound



# Checking the Upper timestamp Bound

